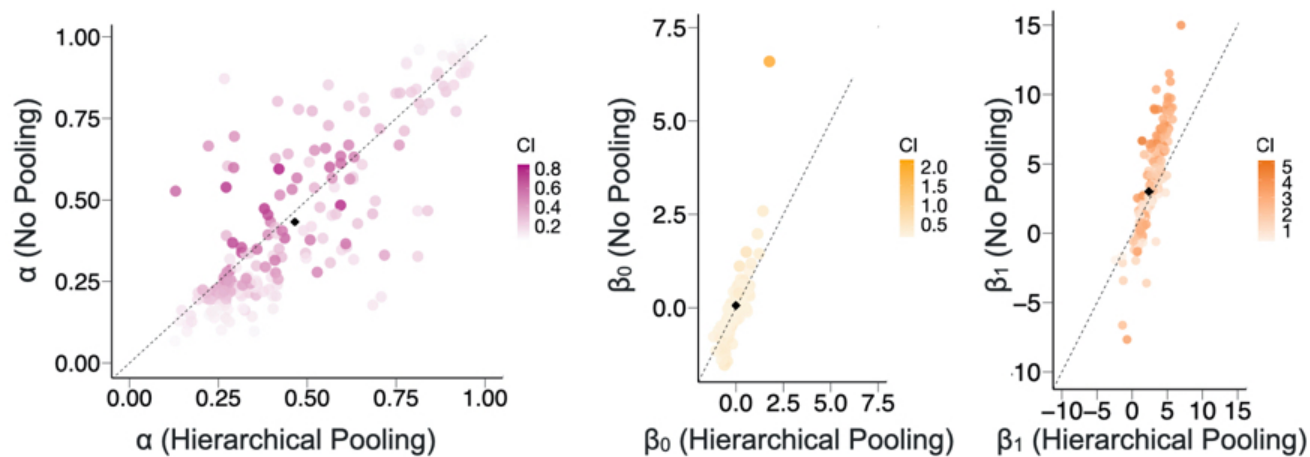


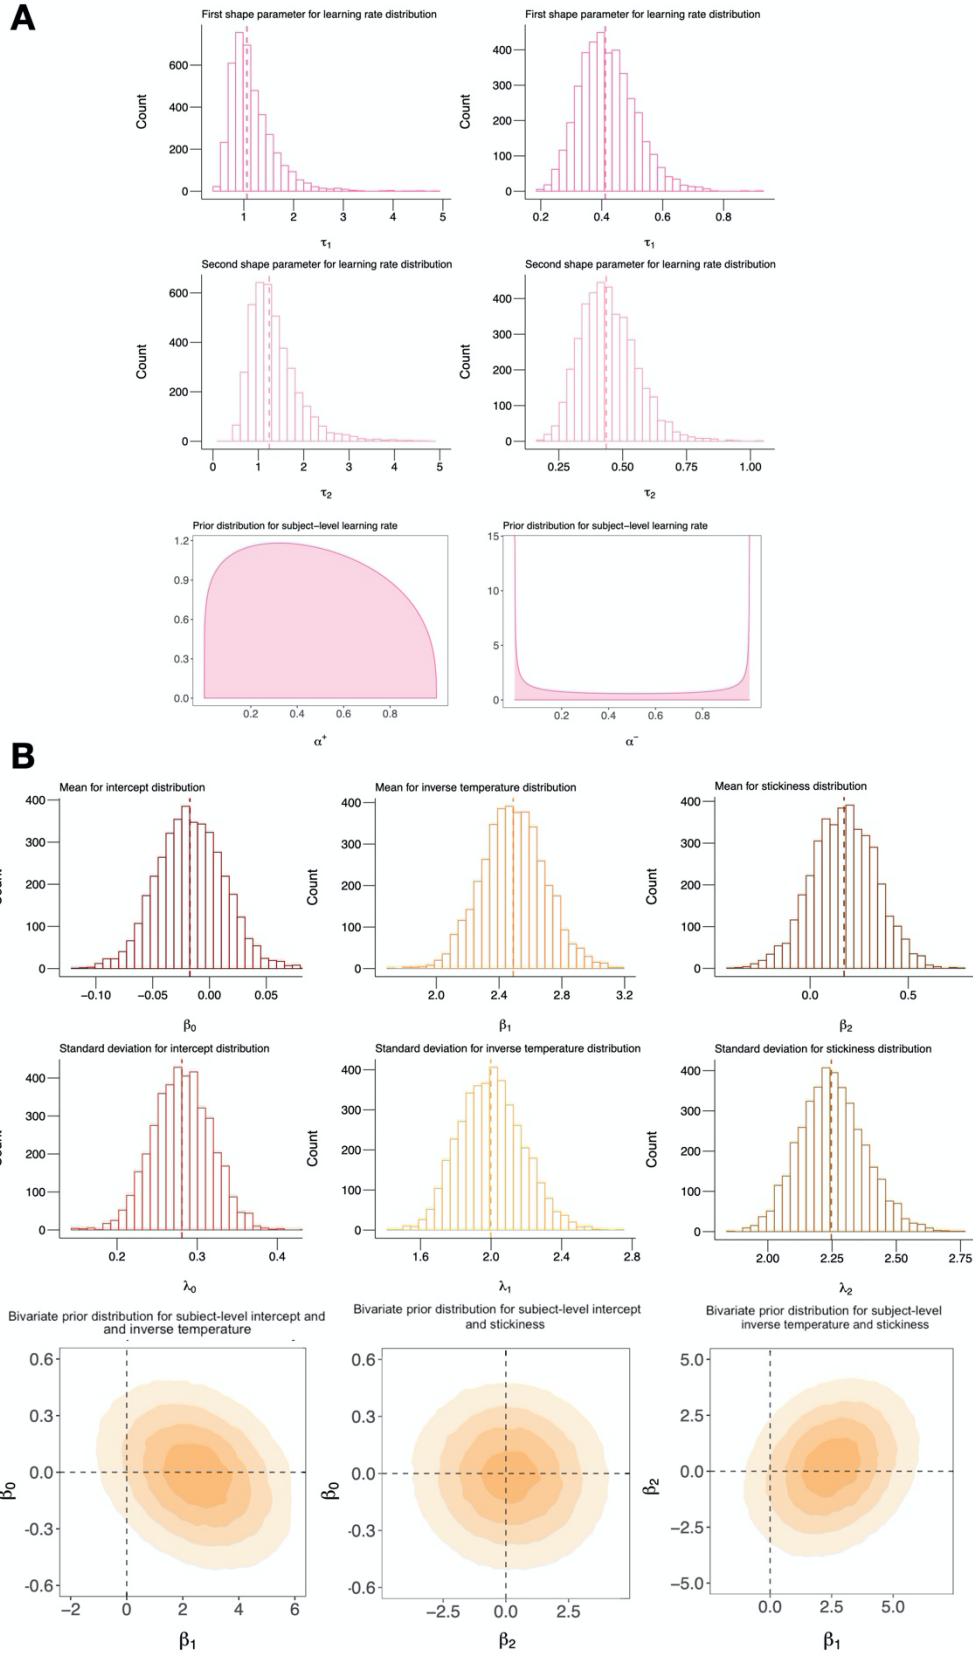
## Supplementary Material

### Hierarchical Bayesian Models of Reinforcement Learning: Introduction and comparison to alternative methods

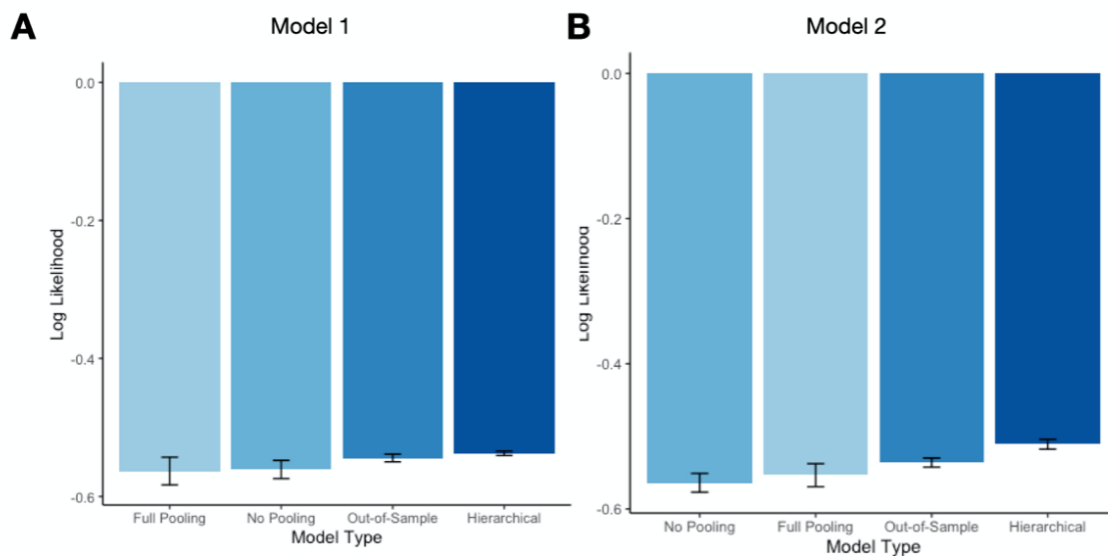
Camilla van Geen and Raphael T. Gerraty



**Figure S1: The use of hierarchical priors helps regularize parameter estimates of initially uncertain or extreme parameters.** Plots of subject-level parameter estimates calculated using the expected value of each subject's posterior distribution from the no pooling vs. hierarchical pooling models. For unconstrained parameters  $\beta_0$  and  $\beta_1$ , hierarchical priors shrink extreme parameter values towards the group mean. The effect is less clear for  $\alpha$ . Shading corresponds to the width of the confidence interval around each subject's parameter estimate in the no pooling model, taken as an index of uncertainty.



**Figure S2: Posterior histograms for group-level parameters and expected empirical prior distributions in model M3. (A) Top:** Posterior histograms for group-level parameters on learning rate estimates, fit hierarchically on model M3. Dashed line corresponds to the median of the each distribution. **Bottom:** Expected empirical prior distribution for learning rates in model M3. While the prior distribution for  $\alpha^+$  provides support for incremental learning — the most likely values of  $\alpha^+$  are between 0.2 and 0.6 — the prior distribution for  $\alpha^-$  highlights that participants may respond to negative outcomes either by ignoring them ( $\alpha^- = 0$ ) or using a win-stay-lose-shift strategy ( $\alpha^- = 1$ ). **(B) Top:** Posterior histograms for group-level parameters on intercept, inverse temperature, and stickiness estimates, fit hierarchically on model M3. **Bottom:** Prior distributions for intercept, inverse temperature, and stickiness, parameterized using the mean of the posteriors. Each plot is multivariate Normal, as it includes model-derived covariance estimates for each combination of parameters ( $\beta_0$  and  $\beta_1$ ;  $\beta_0$  and  $\beta_2$ ;  $\beta_1$  and  $\beta_2$ ).



**Figure S3: Hierarchical pooling improves predictive accuracy for M1 and M2.** Mean of the log-likelihoods for held-out Block 4 data across 40 participants for each of the four candidate model-fitting techniques. Value closer to zero indicate higher predictive accuracy. Error bars reflect within-subject differences based on the method described in Cousineau (2015).